Predicting AI Results on New Benchmarks with a Gaussian Process IRT Model Fit at Test Time

Uriel Nguefack Yefou
Team *Pooled Priors*
uriel.nguefackyefou@mail.concordia.ca

Abstract

We predict whether an AI system answers a benchmark item correctly. The test benchmarks are not in the training data, and labels come a few at a time: 0, 1, 3, 7, 15 and 31 per system and benchmark. For each test benchmark we fit a logistic item response model at test time, using the labels of all systems on that benchmark. Item effects have a Gaussian process prior, so items with close text get close difficulty. We average the model over a grid of prior scales, with weights from the Laplace marginal likelihood. Two things mattered most. First, with few labels most of the error comes from not knowing how well a system does on the benchmark overall, and our prior for this level was centred at the wrong value. Training groups of the size used at test time have a mean accuracy near 0.36, and the test groups we could read from the formative feedback have 0.35. Moving the centre of the prior to this level lowers the Brier score at budget 0 from 0.248 to 0.236 and the Brier ALC by 0.002. This is the largest single gain we found. Second, a linear kernel on item embeddings, centred within each benchmark, is the part of the item kernel that helps most once a benchmark has a few dozen labels. In a simulation of the protocol on held-out benchmarks, the final model has a Brier ALC of 0.1797. A group mean gets 0.1869 and a constant 0.5 gets 0.25.

1 Task and data

A subject is an AI system with fixed settings. An item is one task, with its text, some metadata and an anonymous benchmark ID. A group is one subject on one benchmark. In each scored group with at least 80 items, the items are split in two halves. Labels are shown from one half and the other half is scored. Labels come in nested steps of 0, 1, 3, 7, 15 and 31 per group [AIMS Lab, Stanford University, 2026]. The predictor sees the labels of every group, also of groups it does not predict. The score is the Brier ALC [Brier, 1950], $0.1B_0 + 0.2B_1 + 0.2B_3 + 0.2B_7 + 0.2B_{15} + 0.1B_{31}$, averaged over groups.

We use only the public training pool `aims-foundations/measurement-db` [Truong et al., 2026] (revision `54074453`). We keep benchmarks with binary answers per item and `release = public`, which gives 160 benchmarks. Of these, 92 have at least one group with 80 to 1000 items. This gives 10,208 groups, 154,628 items and 4.3M subject and item cells. Each cell has n trials and k successes. We use no other evaluation data.

2 Simulating the protocol offline

The test benchmarks never appear in training. So we made every choice with a simulator of the official protocol and 5-fold cross-validation over benchmarks.

40th Conference on Neural Information Processing Systems (NeurIPS 2026). Workshop: Predictive AI Evaluation Competition (PAIEC).

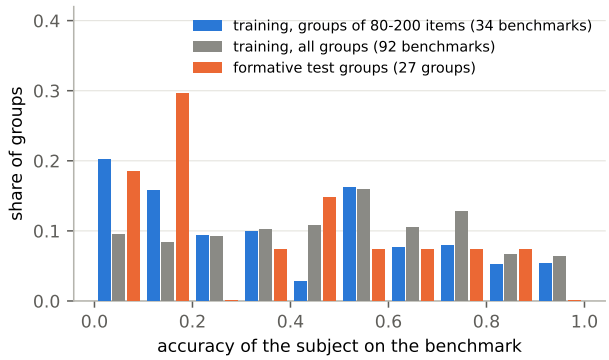


Figure 1: Accuracy of a subject on a benchmark, for training groups and for test groups seen in formative feedback.

Draws. We add groups one at a time. Groups whose benchmark or model name is already in the draw go last, and ties are broken at random. We stop when the draw would go over 1000 subject and item pairs. We use three settings. The realistic setting is our main one. It only allows groups with at most 200 items. A draw then has on average 6.9 groups, 5.9 benchmarks and 142 items per group. The large group setting allows all groups with 80 to 1000 items and gives about 3 groups per draw. In the many subjects setting the limit counts distinct items, so one benchmark can have hundreds of subjects.

Split, labels and scoring. In each group we split the distinct items at random, and the odd item goes to the label pool. Labels come in random order, with one sampled trial per item, and every group sees the labels of all other groups. We compute the exact Brier score of each group over all response rows, apply the ALC weights, and average over groups and then over episodes. We always compare two models on the same episodes, with the same splits and the same labels. We report the mean of the paired differences and its standard error. We chose settings on one seed and checked them again on episodes drawn with another seed.

What the formative feedback tells us. Codabench [Xu et al., 2022] scores each submission on a new draw of about nine groups. The scoring log gives the Brier score and the calibration error (ECE) of every group at every budget. Since the draw changes each time, these logs say little about which model is better. But they do tell us what the test groups look like. A draw has about nine groups from seven benchmarks, with 88 to 152 items per group. This is what our realistic setting gives. The logs also give the accuracy of each group. Our first three submissions predicted exactly 0.5 at budget 0. In that case the ECE of a group is $|q - 0.5|$, where q is its accuracy. Later submissions predicted a little below 0.5 at budget 0, so their Brier score is above 0.25 only when $q > 0.5$, and this tells us the sign. On the 27 groups we recovered this way, the mean accuracy is 0.35. Figure 1 compares them with the training groups. The test groups look like the training groups of realistic size (mean 0.36 per group, 0.39 per benchmark). They do not look like the training pool as a whole (0.49 per benchmark), where large benchmarks with many subjects weigh the most. So the logs told us which setting to use, and that the level of the shift prior was worth testing. We then chose the value of every parameter, including the centre of that prior, by cross-validation on the training pool.

3 Method

3.1 Model

Fix a benchmark b and let $\mathcal{L}$ be the labels shown on b, from any subject. For subject s and item t,

$$\operatorname{logit} P(y_{st} = 1) = c_t + m + a_s + d_t, \quad a_s \sim \mathcal{N}(0, \sigma_a^2),\ d \sim \mathcal{GP}(0, \sigma_d^2 K). \tag{1}$$

Here c_t is a fixed offset from a prior model (Section 3.5), m is a shift for the benchmark, a_s is a shift for the subject, and d_t is the item effect, which is the same for all subjects of the benchmark. This is a logistic item response model [Lord, 1980] with a Gaussian process prior on the item effects

[Rasmussen and Williams, 2006]. Models of this family are also used to evaluate language models on few items [Maia Polo et al., 2024]. Section 3.4 describes the prior on m. The item kernel is

$$K = (1-\rho_\ell)\big[\rho\, e^{(\cos(e_i,e_j)-1)/\ell} + \rho_f\, \kappa_{\text{tok}}(i,j) + (1-\rho-\rho_f)\,\mathbb{1}[i=j]\big] + \rho_\ell\, \frac{\tilde{e}_i^\top \tilde{e}_j}{\|\tilde{e}\|^2}. \qquad (2)$$

e is the normalised `multilingual-e5-small` [Wang et al., 2024] embedding of the item text and metadata. κ_{tok} is the cosine similarity of `key=value` metadata tokens, where we only count tokens found on at least two labeled items. $\tilde{e} = e - \bar{e}_{\mathcal{L}}$ is the embedding minus its mean over the labeled items of the benchmark. The last term is a Bayesian linear regression of item difficulty on the embedding. It is learned for each benchmark from its own labels.

3.2 Inference

We find the posterior mode with Newton's method. The Hessian block of the subject effects is diagonal, so we remove it with a Schur complement. What is left is a block for m and the labeled item effects. Its size is at most the number of distinct labeled items. One Newton step costs $O(|\mathcal{L}| + n_s n_{\mathcal{L}} + n_{\mathcal{L}}^3)$. The Laplace posterior [Rasmussen and Williams, 2006, Chapter 3] gives a mean μ and a variance v of the logit for any target item, labeled or not, through the Gaussian process conditional of d_t given $d_{\mathcal{L}}$. We predict $\sigma\big(\mu/\sqrt{1+\pi\gamma v/8}\big)$. With $\gamma = 1$ this is the usual probit approximation of the predictive integral [MacKay, 1992]. We use $\gamma = 0.35$ while the target subject has at most three labels on the benchmark, and $\gamma = 0.2$ after that. We chose these values in simulation (Section 4). We checked our code against a dense Laplace solver, and the two agree to machine precision.

3.3 Averaging over prior scales

Good prior scales are very different from one benchmark to the next. On some benchmarks items are almost always right or always wrong. On others the item effects are small. So we fit the model for every combination of $\sigma_a \in \{0.4, 0.8, 1.6\}$, $\sigma_d \in \{0.8, 1.6, 3.2\}$ and the three parts of the shift prior below, with $\rho = 0.73$, $\rho_f = 0.17$, $\ell = 0.12$ and $\rho_\ell = 0.4$. We average the predicted probabilities with weights proportional to the Laplace marginal likelihood, as in Bayesian model averaging [Hoeting et al., 1999]. We do this for each benchmark and each set of shown labels on its own.

3.4 Prior on the benchmark shift

The shift m sets how well a subject does overall on a new benchmark. Its prior is a mixture with equal weights,

$$m \sim \tfrac{1}{3}\big[\mathcal{N}(\mu_0, 0.8^2) + \mathcal{N}(\mu_0, 1.6^2) + \mathcal{N}(\mu_0, 3.2^2)\big], \qquad \mu_0 = -0.6, \qquad (3)$$

and each part is one more axis of the averaging grid. The wide part lets the model follow benchmarks where a subject gets almost everything right or almost everything wrong. The narrow part keeps predictions stable with one or two labels. The centre μ_0 matters more than the scales. In training groups of realistic size the mean logit accuracy is -0.87, with a standard deviation of 1.8. So a prior centred at zero starts every new benchmark from an accuracy that is too high. We set μ_0 by cross-validation (Section 4). The value we chose is between zero and the training mean. Two things explain why it does not need to reach -0.87. The offset c_t has a negative mean and already carries part of the low level. And the probit correction pulls predictions toward 0.5.

3.5 Prior across benchmarks and second model

We take hidden states of Qwen2.5-7B-Instruct [Qwen Team et al., 2024] (layer 21 of 28, mean over tokens, standardised) and fit a ridge regression to the mean logit accuracy of each item, with equal weight for all training benchmarks. The output is z_t, and we set $c_t = 0.3\, z_t$. With no labels the prediction is the prior predictive, averaged over the parts of the shift prior. The final predictor is the average of two copies of the model. They differ only in the item vectors used by the kernel: e5-small embeddings in one, Qwen hidden states in the other.

Table 1: Simulated Brier scores on held-out benchmarks (1000 episodes per setting, the same in every row).

	realistic							large	many
Model	B_0	B_1	B_3	B_7	B_{15}	B_{31}	ALC	ALC	ALC
Constant 0.5	.2500	.2500	.2500	.2500	.2500	.2500	.2500	.2500	.2500
Group mean, Beta(1,1)	.2500	.2121	.1862	.1708	.1619	.1572	.1869	.2083	.2133
Bayes, own labels, empirical prior	.2356	.2072	.1856	.1715	.1626	.1576	.1847	.2089	.2129
IRT, no item effects	.2500	.2077	.1835	.1698	.1616	.1571	.1852	.2096	.2047
IRT, independent item effects	.2500	.2113	.1883	.1741	.1652	.1599	.1888	.2091	.1799
GP-IRT, e5 RBF kernel	.2500	.2083	.1846	.1700	.1605	.1540	.1851	.2075	.1738
+ metadata token kernel	.2500	.2090	.1839	.1684	.1583	.1510	.1840	.2071	.1728
+ linear kernel, centred emb.	.2500	.2091	.1856	.1688	.1582	.1507	.1844	.2046	.1726
+ Qwen offset c_t	.2483	.2082	.1851	.1685	.1580	.1507	.1839	.2041	.1729
+ second model, Qwen kernel	.2483	.2083	.1852	.1684	.1578	.1502	.1838	.2039	.1709
+ retuned scales and γ	.2483	.2055	.1829	.1672	.1571	.1497	.1824	.2038	.1701
+ shift scale mixture, γ by labels	.2483	.2046	.1822	.1668	.1568	.1494	.1819	.2046	.1698
+ shift prior centred at -0.6 (final)	**.2356**	**.2012**	**.1814**	**.1666**	**.1567**	**.1494**	**.1797**	.2041	**.1693**

3.6 Label choice and run time

We keep the platform's default random choice of labels (Section 5). Each prediction worker keeps one fit per benchmark and set of labels in a cache. The label list only grows, so we can tell that a set has not changed from its length and its last entry. Item vectors are cached too. Without a GPU, or if Qwen cannot be loaded, the predictor uses the e5 model alone with no offset c_t. If Qwen fails during the run, we turn it off and go on with the e5 model. In simulation the e5 model alone is worse by only 0.0003 ALC in the realistic setting. A prediction takes 15 to 110 ms on the A100 we used for testing.

4 Results

Table 1 adds the parts of the model one at a time, on the same 1000 episodes in every row. In the realistic setting, the paired standard error of the difference between two rows is about 10^{-4}. The centred shift prior gives the largest step in the realistic setting (-0.0022, from 0.1819 to 0.1797), and almost all of it comes at budgets 0 and 1. The centred linear kernel gives the largest step with large groups (-0.0025) and helps from budget 7 on. The final model is the best row in the realistic and many subjects settings. With large groups it is within 0.0003 of the best row, which is inside the noise of that setting. Most of the other rows change the ALC by less than 0.001.

Where the error comes from. Two oracles show how far a predictor can go in the realistic setting. If we knew the accuracy of each group, a constant prediction per group would score 0.1503. If we knew the accuracy of each item for each subject, the Brier score would be 0.011. Our B_{31} of 0.1494 is just below the first oracle. With 31 labels the model knows the group level well and gets a bit more from the items. But the gap to the second oracle shows that most of the item signal is still out of reach. The row *Bayes, own labels* in Table 1 is a more useful point of comparison. It is the Bayes predictor of the group accuracy from the group's own labels. Its prior is the accuracy of training groups from other folds (groups of 80 to 200 items, in all three columns). At budget 0 it predicts the mean training accuracy, which is the best constant under the training distribution. Its B_0 is 0.2356, and the final model gets the same value. So without information on single items, there is little left to gain at budget 0. From budget 1 on, the final model beats this reference by 0.004 to 0.008 at every budget. This part of the gain comes from the labels of other subjects on the same benchmark and from the item kernel. Figure 2 shows the curves.

Prior on the benchmark shift. We compared centres and scales on 300 paired episodes of the realistic setting, against the model of the row before it in Table 1. Centres of -0.4, -0.6 and -0.8 lower the ALC by 0.0018 to 0.0022, all with a standard error near 0.0002, while -1.2 gives only 0.0010. The scales matter less: $(1, 2, 4)$, $(0.8, 1.6, 3.2)$ and $(1.5, 3)$ are within 0.0001 of each other. Gaussian mixtures with two or three parts, fit directly to the accuracies of training groups, did about as well (-0.0023), so we kept the simpler prior. A larger weight of 0.5 for the Qwen offset did

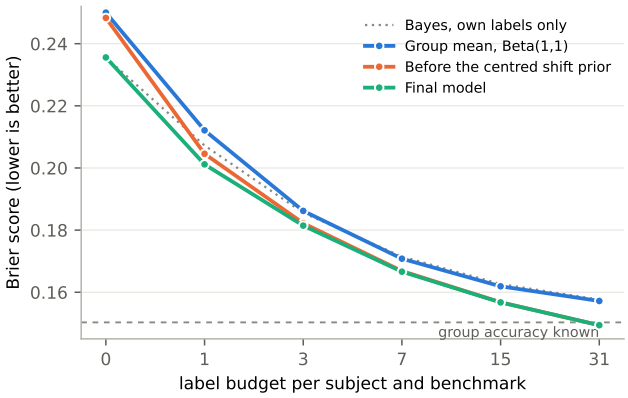


Figure 2: Brier score by budget in the realistic setting.

not add to the gain, and neither did more confident predictions with few labels. Once the prior is centred, the Qwen offset no longer matters: setting $c_t = 0$ changes the ALC by $+0.00001 \pm 0.00007$. We kept it anyway, since we had tested the full runtime with it. On 1000 new episodes drawn with another seed, the final model beats the row before it by 0.0022 ± 0.0001 ALC. The Brier score drops by 0.012 at budget 0 and by 0.004 at budget 1. In the large group setting, where training accuracy is higher, the centred prior neither helps nor hurts. The change is -0.0008 ± 0.0005 on the episodes used for the choice and $+0.0002 \pm 0.0003$ on new episodes.

Risk with very few labels. With one to three labels, a predictor can move away from the truth. In simulation, $B_3 > B_0$ happens for about a quarter of the groups. In the 53 groups of our first formative logs it happened for 24.5%, so the simulator matches what we saw on the platform. A smaller variance scale γ makes predictions more confident. It lowers the mean ALC down to about $\gamma = 0.1$, but it also makes bad groups more common: the share of groups with $B_3 > 0.3$ goes from 5.2% at $\gamma = 0.5$ to 7.4% at $\gamma = 0.2$. We kept $\gamma = 0.35$ while the subject has at most three labels, and $\gamma = 0.2$ after that, when the posterior knows more.

Formative feedback. The final model (submission 1018959) scored 0.182 Brier ALC on Codabench, with an ECE of 0.139. Our seven earlier submissions scored between 0.164 and 0.228. Each submission is scored on a new draw of eight or nine groups, so these numbers move by several hundredths from one draw to the next, and we did not use them to rank models.

5 What did not help

We tested each of these ideas with the same paired simulation and dropped them.

- Predicting budget 0 from item text and subject attributes with LightGBM [Ke et al., 2017] (e5 embeddings, text statistics, provider, release date, size, model name). The best B_0 was 0.248, against 0.250. Once the shift prior is centred, the model already reaches the B_0 of the Bayes reference (0.2356).
- Asking Qwen2.5-7B to rate item difficulty from 1 to 9. Inside a benchmark, the rank correlation with measured difficulty was -0.03. Items rarely come with a reference answer (3 benchmarks out of 92), so we could not score the likelihood of the correct answer either.
- Looking for the same item in other benchmarks. Fewer than 0.1% of items have a match with cosine above 0.95.
- Learning a projection of the embeddings for the kernel by maximising the marginal likelihood on training groups. It was worse on held-out benchmarks. A learned low-rank metric for the linear kernel did not help either.
- Bigger item vectors (e5-base, e5-large, Qwen2.5-7B layers 14, 21 or 28) as the only kernel input. None beat e5-small. With so few labels, the model is limited by the number of labels

more than by the quality of the vectors. Using Qwen2.5-14B instead of the 7B model gained 0.0002 ALC and needed twice the memory.

- Stacking with LightGBM on model outputs and label counts, and Platt scaling [Platt, 2000] per budget. Both were worse on held-out benchmarks, because calibration learned on some benchmarks does not carry over to others.
- Other ways to choose labels: uncertainty, overlap with other subjects, embedding diversity, D-optimal design. None was clearly different from random order.
- Learning the centre of the shift prior during the run, from the label rates of the other benchmarks in the same draw (a hierarchical prior with a shared centre). It was worse in every fold, by 0.0001 to 0.0007 ALC depending on the prior width of the centre, even in the fold with the hardest benchmarks. How well one subject does on one benchmark says little about how other subjects do on other benchmarks.
- Weights for the averaging grid learned on training benchmarks. The posterior over the grid stays almost flat even with 31 labels (weights from 0.026 to 0.048 over 27 points), so a learned prior is close to the uniform one.
- A prior on subject ability from the same model name on other benchmarks. Its out-of-fold correlation with ability on held-out benchmarks was -0.16, and it gave no gain. A ridge regression of the group level on all subject metadata tokens did no better. Its out-of-fold correlation with the group logit was between -0.19 and $+0.24$ depending on the subset, and on all groups it was worse than a constant. A prior on item difficulty inside a benchmark, from a ridge on centred Qwen states (Spearman 0.07), hurt B_0 and B_1 more than it helped later.
- Numeric metadata (lengths, counts and similar fields) or text statistics as an extra kernel term. All versions were 0.0001 to 0.0005 worse.

6 Limitations

The simulator depends on how we read the draw rule. We matched it to the formative logs, and for this reason we checked every choice in three settings. The centred shift prior helps because test groups have a low accuracy on average, like the training groups of the same size. On a test set made mostly of easy benchmarks, it would lose part of its gain at budget 0, although it did not hurt in our large group setting. The Qwen parts need a GPU at test time. Without one, the predictor still runs with the e5 model alone. The model assumes that item effects are the same for all subjects of a benchmark, apart from the logistic link. Last, with no labels the model predicts close to the average level of a group, because the data we have says little about how well a subject will do on a new benchmark.

7 Reproducibility

Code: https://github.com/nguefackuriel/paiec-gpirt. The README goes through the full pipeline: download of the training pool, the tables, the benchmark folds, the embeddings, the simulator (`paiec/sim.py`), every experiment script (`scripts/try_*.py`, `scripts/final_table.py`), the Qwen prior (`scripts/llm_embed.py`, `scripts/build_qwen_prior.py`), the check that the submission gives the same predictions as the simulator (`scripts/parity_v8.py`), and packaging. The script `scripts/formative_acc.py` recovers the accuracy of test groups from the formative logs. Pretrained models: `intfloat/multilingual-e5-small` and `Qwen/Qwen2.5-7B-Instruct`. Training data: only `aims-foundations/measurement-db`. Compute: one A100 80GB for the hidden states (about 2.5 hours) and 30 CPU cores for the simulations. One setting of 1000 episodes takes 10 to 20 minutes in the realistic setting.

Use of AI assistants. We wrote the code, ran the experiments and drafted this report with help from an LLM coding assistant (Claude Code, Anthropic). All results come from the code in the repository.

References

AIMS Lab, Stanford University. Predictive AI evaluation competition (PAIEC) at NeurIPS 2026. https://aims.stanford.edu/competition, 2026. Accessed 30 October 2026.

Glenn W. Brier. Verification of forecasts expressed in terms of probability. *Monthly Weather Review*, 78(1):1–3, 1950. doi: 10.1175/1520-0493(1950)078<0001:VOFEIT>2.0.CO;2.

Jennifer A. Hoeting, David Madigan, Adrian E. Raftery, and Chris T. Volinsky. Bayesian model averaging: A tutorial. *Statistical Science*, 14(4):382–417, 1999. doi: 10.1214/ss/1009212519.

Guolin Ke, Qi Meng, Thomas Finley, Taifeng Wang, Wei Chen, Weidong Ma, Qiwei Ye, and Tie-Yan Liu. LightGBM: A highly efficient gradient boosting decision tree. In *Advances in Neural Information Processing Systems*, volume 30, 2017.

Frederic M. Lord. *Applications of Item Response Theory to Practical Testing Problems*. Lawrence Erlbaum Associates, Hillsdale, NJ, 1980.

David J. C. MacKay. The evidence framework applied to classification networks. *Neural Computation*, 4(5): 720–736, 1992. doi: 10.1162/neco.1992.4.5.720.

Felipe Maia Polo, Lucas Weber, Leshem Choshen, Yuekai Sun, Gongjun Xu, and Mikhail Yurochkin. tinyBenchmarks: Evaluating LLMs with fewer examples. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 34303–34326, 2024.

John C. Platt. Probabilities for SV machines. In *Advances in Large-Margin Classifiers*, pages 61–74. The MIT Press, 2000. doi: 10.7551/mitpress/1113.003.0008.

Qwen Team, An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, et al. Qwen2.5 technical report. *arXiv preprint arXiv:2412.15115*, 2024.

Carl Edward Rasmussen and Christopher K. I. Williams. *Gaussian Processes for Machine Learning*. The MIT Press, 2006. doi: 10.7551/mitpress/3206.001.0001.

Nhi Truong, Sang T. Truong, and Sanmi Koyejo. The AI measurement data bank. https://aimslab.stanford.edu/measurement-db, 2026. AIMS Lab, Stanford University.

Liang Wang, Nan Yang, Xiaolong Huang, Linjun Yang, Rangan Majumder, and Furu Wei. Multilingual E5 text embeddings: A technical report. *arXiv preprint arXiv:2402.05672*, 2024.

Zhen Xu, Sergio Escalera, Adrien Pavão, Magali Richard, Wei-Wei Tu, Quanming Yao, Huan Zhao, and Isabelle Guyon. Codabench: Flexible, easy-to-use, and reproducible meta-benchmark platform. *Patterns*, 3(7):100543, 2022. doi: 10.1016/j.patter.2022.100543.